

CS4375 Operating Systems Concepts
Fall 2022. Homework Assignment 1
Due: 09/21/2021 11:59PM MDT
Group assignment

This first assignment is a basic Linux get-to-know that will provide the necessary basis for the following assignments, as well as some basic programs manipulating files, written in C, as training to get everyone up to speed with this programming language. The assignment is to be done in a group of three (3) students. We have already formed the teams in class. **There is quite some amount of code to be written, start with a meeting now, distribute the tasks, and start coding this week!**

Even though all homework assignments in this class will be group assignments, the instructor and grader will determine each student's **individual** involvement. **Your grade will depend on your individual performance.**

For this assignment, you have to turn in an archive containing:

- the script `findlocation`,
- the source file `head.c`,
- the source file `tail.c`,
- the source file `findlocationfast.c` and
- your report (as a PDF) covering all tasks in all sections of this assignment.

The Linux manual pages, which you can open using the shell command e.g. `man 2 mmap`, are your friend. Read and understand them in their entirety for all system calls that are used in the programs to be written in this assignment.

Even though it is a call to a level-3 function, your programs may use calls to `strerror` for error messages. None of the programs is supposed to use `printf`, `fprintf` or `fputs`, as these are level-3 calls. You might want to start with the implementation of a function `int my_file_puts(int fd, const char *s);` which you base on the `write` system call (and your implementation of a function performing the task of `strlen`).

1 Linux ssh Login

This first part of the assignment is really easy: you are just required to access the machine `titanic` at `titanic.christoph-lauter.org` through `ssh` on **port 2233**. Using `ssh` on some Linux/MacOS/FreeBSD/whatever machine or Putty on Windows and using the login and password that have been announced in class, access the machine and change your password to something other than the default value. If you wish, put a public `ssh` key onto `titanic` so that `ssh` will never ask again for the password (once you changed it). The following rules govern the use of the `titanic` system:

- You need to change your password from the default password to a new one as soon as possible.
- **All accounts that still have the default password on 09/10/2022 will be deleted and will not be recreated.**
- You must force the port used for the `ssh` connection to port 2233 (and not on the regular port 22).
- Another system answers on port 22 (the default `ssh` port). You have **no** account on that other system. **Hammering the other system results in getting banned (by `fail2ban`) on all subdomains `*.christoph-lauter.org`. Make sure to use the right port 2233.**
- The system runs `fail2ban`. If you fail to authenticate correctly too many times, your IP will get banned for 24 hours. If you need to get unbanned, send email (urgent cases only).

- It is recommended that you put a private `ssh` key onto the system to avoid password authentication.
- Unless you turn on the `VerifyHostKeyDNS` option on in `ssh` (set it to yes), the first time you connect to the server, you need to accept its fingerprint. Using the `VerifyHostKeyDNS` option is recommended but not all versions of `ssh` support it.
- This system runs in my **private** subnet at **home**. I **personnally** pay for the electricity consumed by the system.
- **Any attempt at exploiting the system for the purpose of sending spam, mining bitcoins or “just” hacking my personal network or any other purpose which is not related to this class will result in a notice to UTEP and other consequences, which might include you failing this class!**
- If you f**k up on the system (have a process that takes all memory, uses all CPUs, writes to disc like crazy etc.) but it’s just that you happen to do that inadvertently, **no biggie! Just send email as soon as possible.**
- **Never take the system’s resources for granted.** If the computer needs rebooting or there’s a power failure from El Paso Power, your data/processes etc. may just go away. This should never happen, but who knows. **Make sure enough time is left for homework submissions. Make sure to backup your important data that is on that server.**
- You can use `scp` to move data from that server resp. to that server.

2 Getting to know the `bash` shell

Open a terminal and perform the following on command line. Keep a log of your commands and the answers you got back; this will provide you a basis for the write-up you turn in for this assignment.

1. Using `ssh`, log into the remote system `titanic` on `titanic.christoph-lauter.org`.
2. On that system find out
 - how many users are logged in concurrently with you,
 - how long the system is already running since last rebooted,
 - what processors it has,
 - how much RAM is available.
3. Copy the file `nanpa`, found in the home directory of the `cqlauter` user, to your account or use `scp` to copy the file back home onto your private computer.
4. On `titanic` or a private Linux installation, open the file `nanpa` using `less`. The file contains quite a comprehensive list of North American phone number prefixes (first 6 digits, excluding +1), followed by the location this phone number prefix is attached to. For example, for 915220, the location El Paso TX is listed. Still inside `less`, find the entries for 907519, 503526 and a couple of other phone numbers you know in the country, as such your home phone, your parents’ phone, the phone of a loved one etc.
5. Find out how many lines connecting prefixes to locations are contained in the file `nanpa`. Which Linux command line tool do you use to count lines?
6. List the first 17 lines of the `nanpa` file on command line. Also list the last 42 lines of the file. You can use the Linux tools `head` and `tail` for this task.

7. Write a short `bash` script `findlocation` that takes a 6-digit prefix in argument and displays the corresponding location. If the script receives no argument or the argument is not a 6-digit prefix made only out of the digits 0 through 9, the script must return an exit condition code signaling failure (e.g. by executing `exit 1`). If the script receives a correctly formatted argument but the prefix is not found in the `nanpa` file, the script must return an exit condition code signaling failure. Otherwise, the script must display the appropriate location (on `stdout`). The location must not be prefixed by the prefix nor followed by superfluous spaces. This means you have to format the line found in the `nanpa` file before displaying it. You may use `grep` and `sed` for this script.

3 head and tail

As seen above, the Linux/UNIX/POSIX `head` and `tail` programs allow the beginning resp. the end of a text file to be extracted.

In this section of the homework assignment, you are asked to re-program the Linux `head` and `tail` commands. Your programs must be written in C and must be based solely on POSIX file-system handling system calls (level-2 calls). This means, you can use `open`, `close`, `read` and `write` but calls to `fopen`, `fread` or `fprintf` are prohibited. You may use `malloc`, `calloc` and `free`; however, see the remark below.

Your `head` and `tail` programs must handle the following use-cases and associated options:

- When run without any option nor filename argument, the `head` command must copy the first 10 lines of the input it receives on standard input (`stdin`) to standard output (`stdout`) and disregard any other lines seen on input. Similarly, the `tail` command must read the input on standard input and output only the last 10 lines of that input to standard output. Note that `head` can perform its task without any buffering (in memory), while `tail` needs to maintain a memory buffer of the current last 10 lines it just read in. When the input contains less than 10 lines, the input is completely copied from standard input to standard output.
- When run with a filename argument, the `head` and `tail` commands do not read from standard input but from the file designated by that filename argument. If they cannot open the file in read-mode, they display an appropriate error message on standard error and exit. Your `head` and `tail` commands just handle one filename argument, whereas the Linux/UNIX/POSIX programs handle several filename options.
- When run with the `-n [num]` option, the `head` and `tail` commands behave as specified above but replace the number of 10 lines by the appropriate amount specified by `[num]`. If that amount is zero, they do nothing. If the amount is negative, your `head` and `tail` programs do nothing either; the Linux/UNIX/POSIX commands execute a special behavior in this case. When the `-n` option is given incorrectly or incompletely, e.g. by specifying `-nose` instead of `-n` or by specifying `-n` without specifying a non-negative amount in the next argument, the `head` and `tail` commands display an appropriate error message on standard error and exit. The amount `[num]` is supposed to be less than or equal to $2^{64} - 1$; if it is larger, the behavior of the programs is unspecified. You are supposed to perform the string-to-integer conversion with code written on your own; not using `atoi` or `strtol`, which would be level-3 operations.
- The filename and `-n` options can be combined and used in any reasonable order, i.e. as `-n [num] [filename]` or as `[filename] -n [num]`.
- If no error condition arises, the `head` and `tail` commands return the *success* condition code, i.e. zero. If an error occurs, they return a non-zero *failure* condition code. Note that an error might occur for any system call, i.e. for any call to `read`, `write`, `malloc` etc. Before exiting, an error message must be displayed on standard error in these cases, too.

These are examples of well-formed calls to `head` and `tail`:

```
cat nanpa | ./head
cat nanpa | ./head -n 42
./head -n 42 nanpa
./head nanpa
./head nanpa -n 42
echo -e 'Hello\nworld!\nHave\nna\nnice\nday.' | ./head
cat nanpa | ./tail
cat nanpa | ./tail -n 42
./tail -n 42 nanpa
./tail nanpa
./tail nanpa -n 42
echo -e 'Hello\nworld!\nHave\nna\nnice\nday.' | ./tail
```

Your `head` and `tail` commands must not access any memory out of their addressing space and must not leak any memory. This means any memory allocated on the heap with `malloc` or `calloc` must be returned using `free`. All accesses to arrays must use correct indices; you may want to double-check on this. All files opened with `open` must be closed with `close`. These memory and file handling specifications must be met even though the operating system would perform these tasks automatically on process termination if they are forgotten. Your instructor will check for these points, using, amongst others, `strace` and `valgrind`. You may want to utilize these two tools for testing, too.

You must write C code (not C++). You are supposed to do some software engineering in order to end up with well structured source code (no spaghetti!). Your source code must contain some comments at the most decisive places, i.e. before functions to document their meaning and at places where the functioning of the code is not obvious to someone reading your code for the first time.

As the `head` and `tail` commands exist on the system you will be developing on, be sure to execute **your** `head` and `tail` implementations when testing!

For this section, you have to submit:

- The source file `head.c` which can be compiled to `head` using `gcc -Wall -O3 -o head head.c`.
- The source file `tail.c` which can be compiled to `tail` using `gcc -Wall -O3 -o tail tail.c`.
- A section in your report explaining the software engineering decisions you made, the problems you encountered and the testing you performed.

4 File searching

In the Linux shell get-to-know part above, you already worked with the file called `nanpa`. This file contains quite a comprehensive list of North American phone number prefixes, six digits long each, followed by a string describing the location served by phone numbers starting with this prefix. Again, on each line, the string is right next to the six digit prefix. It is suffixed with spaces so it is always exactly 25 digits long. Each line is separated from the next line by one next-line character `'\n'`. Hence each line consists of exactly $6 + 25 + 1 = 32$ characters. The `nanpa` file is sorted by ascending prefixes. However, there are certain prefixes that do not correspond to any location: these prefixes do not figure in the file.

With the script `findlocation` you'll write for the section above, you'll be able to search the `nanpa` file in $\mathcal{O}(n)$ using the Linux/UNIX/POSIX tools `grep` and `sed`. However, this method is not optimal from a complexity standpoint: a look-up in an ordered array or file can be done in $\mathcal{O}(\log n)$.

A typical POSIX-compatible environment supports the `mmap` system call. This system call allows the contents of a file on the filesystem, described by a file handle obtained with `open`, to be mapped into memory: the `mmap` call

returns a pointer which, when dereferenced, allows for reading the memory-mapped file by reading at the address given by the pointer, and for writing to the file, by writing to memory at the address provided by the pointer.

For this section of the homework assignment, you must write, in C, a tool named `findlocationfast` which performs the same task as the script `findlocation`. The difference is that `findlocationfast` uses `mmap` to map the used file into memory.

You are supposed to base your program only on the system calls `open`, `lseek`, `mmap`, `munmap` and `close`. You are not supposed to use level-3 buffered file manipulation calls, like `fopen`, `fread`, `fprintf` etc, not even for error messages. In case of an error, display `An error occurred while xxx.`¹ on standard error which just a call to write. Write your own `display_error_message` function for that purpose. Your tool may assume that the file to search is less than 2147483648 bytes in size. Your tool must not leak any memory, the memory-mapped region must be properly unmapped using `munmap`. All files opened must be properly closed. The status codes to be returned are the same as for `findlocation`. Your tool is not supposed to allocate any memory (using `malloc` or `calloc`) other than the memory obtained with `mmap`; using statically sized buffers on the stack is fine. The tool is not supposed to use `strlen`, `strcmp`, `memset`, `memcpy` or `memmove`; these functions are trivial to implement: just implement your own. The lookup algorithm must have $\mathcal{O}(\log n)$ time complexity. The tool may gracefully² fail on files that are non-regular or non-seekable³. If you choose to implement an additional mode for non-seekable files, document this point; extra points may be given. You must achieve $\mathcal{O}(\log n)$ time complexity for seekable files though⁴.

For this section, you have to submit:

- The source file `findlocationfast.c` which can be compiled to `findlocationfast` using `gcc -Wall -O3 -o findlocationfast findlocationfast.c`.
- A section in your report explaining the software engineering decisions you made, the problems you encountered, the testing you performed and the explanations you found.

¹Where `xxx` stands for the operation you just attempted to execute.

²with a proper error message (on standard error) and deallocating memory and closing all files.

³A file is seekable if the `lseek` system call works on that file. Some files in the Unix/Linux world are non-seekable.

⁴Achieving $\mathcal{O}(\log n)$ time complexity is difficult or even impossible on non-seekable files.